

Data Science II Homework 2

Camille Okonkwo

Contents

- 1a) Fit smoothing spline models to predict out-of-state tuition (Outstate) using the percentage of alumni who donate (perc.alumni) as the only predictor, across a range of degrees of freedom. Plot the model fits for each degree of freedom. Describe the observed patterns that emerge with varying degrees of freedom. Select an appropriate degree of freedom for the model and plot this optimal fit. Explain the criteria you used to determine the best choice of degree of freedom. 4
- 1b) Train a multivariate adaptive regression spline (MARS) model to predict the response variable. Report the regression function. Present the partial dependence plot of an arbitrary predictor in your model. Report the test error. 7
- 1c) Construct a generalized additive model (GAM) to predict the response variable. Does your GAM model include all the predictors? For the nonlinear terms included in your model, generate plots to visualize these relationships and discuss your observations. Report the test error. 10
- 1d) In this dataset, would you favor a MARS model over a linear model for predicting out-of-state tuition? If so, why? More broadly, in general applications, do you consider a MARS model to be superior to a linear model? Please share your reasoning. 27

```
library(tidymodels)
library(splines)
library(caret)
```

Partition the dataset into two parts: training data (80%) and test data (20%) with `tidymodels`.

```
college = read_csv("data/College.csv") |>
  drop_na() |>
  select(-College)

set.seed(2)

# create a random split of 80% training and 20% test data
data_split <- initial_split(data = college, prop = 0.8)

# partitioned datasets
training_data = training(data_split)
testing_data = testing(data_split)

head(training_data)

## # A tibble: 6 x 17
##   Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad Outstate
##   <dbl> <dbl> <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1  1380   768   263      57      82     1000     105   19300
## 2   434   321   141      28     53      624     269   10950
## 3  2013  1053   212      33     61      912     158   5150
## 4  2324  1319   370      52     81     1686      35  16560
## 5  1709  1385   634      36     72     2281      50  14125
## 6   427   385   143      18     38      581     533  12700
## # i 9 more variables: Room.Board <dbl>, Books <dbl>, Personal <dbl>, PhD <dbl>,
## #   Terminal <dbl>, S.F.Ratio <dbl>, perc.alumni <dbl>, Expend <dbl>,
## #   Grad.Rate <dbl>
```

```
head(testing_data)
```

```
## # A tibble: 6 x 17
##   Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad Outstate
##   <dbl> <dbl> <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1  2186  1924   512      16     29     2683     1227  12280
## 2  1428  1097   336      22     50     1036      99  11250
## 3   193   146    55      16     44      249     869   7560
## 4   582   498   172      21     44      799      78  10468
## 5  1732  1425   472      37     75     1830     110  16548
## 6   494   313   157      23     46     1317     1235  8352
## # i 9 more variables: Room.Board <dbl>, Books <dbl>, Personal <dbl>, PhD <dbl>,
## #   Terminal <dbl>, S.F.Ratio <dbl>, perc.alumni <dbl>, Expend <dbl>,
## #   Grad.Rate <dbl>
```

```
# training data
x <- model.matrix(Outstate ~ ., training_data)[, -1] # matrix of predictors
head(x)
```

##	Apps	Accept	Enroll	Top10perc	Top25perc	F.Undergrad	P.Undergrad	Room.Board
## 1	1380	768	263	57	82	1000	105	6694
## 2	434	321	141	28	53	624	269	4600
## 3	2013	1053	212	33	61	912	158	3036
## 4	2324	1319	370	52	81	1686	35	5140
## 5	1709	1385	634	36	72	2281	50	3600
## 6	427	385	143	18	38	581	533	5800

##	Books	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend	Grad.Rate
## 1	600	700	89	93	6.1	18	14779	83
## 2	550	950	79	82	12.9	30	9264	81
## 3	500	1655	64	74	10.5	11	7547	59
## 4	558	1152	91	93	10.5	30	16196	79
## 5	400	700	79	89	12.5	58	9907	80
## 6	450	700	81	85	10.3	37	11758	84

```

y <- training_data$Outstate # vector of response

# testing data
x2 <- model.matrix(Outstate ~ .,testing_data)[, -1] # matrix of predictors
y2 <- testing_data$Outstate # vector of response

```

1a) Fit smoothing spline models to predict out-of-state tuition (Outstate) using the percentage of alumni who donate (perc.alumni) as the only predictor, across a range of degrees of freedom. Plot the model fits for each degree of freedom. Describe the observed patterns that emerge with varying degrees of freedom. Select an appropriate degree of freedom for the model and plot this optimal fit. Explain the criteria you used to determine the best choice of degree of freedom.


```
# create a grid for x
perc.alumni.grid <- seq(0, max(college$perc.alumni) + 5, by = 1)

# loop prep
fit.ss = list()
pred.ss = list()
pred.ss.df = list()
pred.ss.df.range = data.frame()

set.seed(2)
# loop for a range of degrees of freedom
for (i in 1:20) {
  fit.ss[[i]] = smooth.spline(training_data$perc.alumni, training_data$Outstate, df = i)
  pred.ss[[i]] = predict(fit.ss[[i]], x = perc.alumni.grid)
  pred.ss.df[[i]] = data.frame(pred = pred.ss[[i]]$y, perc.alumni = perc.alumni.grid, df = i)
  pred.ss.df.range = rbind(pred.ss.df[[i]], pred.ss.df.range)
}

# scatter plot
p <- ggplot(data = training_data, aes(x = perc.alumni, y = Outstate)) + geom_point(color = rgb(0.2, 0.4, 0.6))

# plot the model fits for each degree of freedom
p +
  geom_line(aes(x = perc.alumni, y = pred, group = df, color = df), data = pred.ss.df.range) + theme_bw()
```



```
set.seed(2)
```

```
# select an appropriate degree of freedom for the model
fit.ss.optimal = smooth.spline(training_data$perc.alumni, training_data$Outstate)
```

```
fit.ss.optimal$df
```

```
## [1] 2.000245
```

```
# predicted values
```


```
pred.ss.optimal <- predict(fit.ss.optimal,
                           x = perc.alumni.grid)
```

```
pred.ss.optimal.df <- data.frame(pred = pred.ss.optimal$y,
                                 perc.alumni = perc.alumni.grid)
```

```
# plot this optimal fit
```

```
p.optimal <- ggplot(data = training_data, aes(x = perc.alumni, y = Outstate)) + geom_point(color = rgb(0.8, 0.1, 0.1, 1))
```

```
p.optimal +
  geom_line(aes(x = perc.alumni, y = pred), data = pred.ss.optimal.df,
            color = rgb(0.8, 0.1, 0.1, 1)) + theme_bw()
```



When the degrees of freedom is smaller, the model resembles a linear model. As the degrees of freedom increase, the model becomes more flexible and we can see that exemplified through the wavy lines. For the optimal smoothing splines model, the degrees of freedom = 2.0002451. To determine the best choice of degrees of freedom, we can use cross-validation to choose the degrees of freedom that result in the best predictive performance.

1b) Train a multivariate adaptive regression spline (MARS) model to predict the response variable. Report the regression function. Present the partial dependence plot of an arbitrary predictor in your model. Report the test error.


```
# 10-fold cross-validation
ctrl <- trainControl(method = "cv", number = 10)

# set grid
mars_grid <- expand.grid(degree = 1:4, nprune = 1:16)

set.seed(2)

# fit a MARS model
mars.fit <- train(x, y,
                  method = "earth",
                  tuneGrid = mars_grid,
                  trControl = ctrl)

# plot
ggplot(mars.fit)
```



```
# best tuning parameters
```

```
mars.fit$bestTune
```

```
##      nprune degree
```

```
## 12      12      1
```

```
# regression function
```

```
mars.fit$finalModel
```

```
## Selected 12 of 22 terms, and 9 of 16 predictors (nprune=12)
```

```
## Termination condition: RSq changed by less than 0.001 at 22 terms
```

```
## Importance: Expend, Grad.Rate, Room.Board, Accept, Enroll, F.Undergrad, ...
```

```
## Number of terms at each degree of interaction: 1 11 (additive model)
```

```
## GCV 2729781    RSS 1111485831    GRSq 0.8134662    RSq 0.8312207
```

```
# report the regression function
```

```
summary(mars.fit)
```

```
## Call: earth(x=matrix[452,16], y=c(19300,10950,5...), keepxy=TRUE, degree=1,
```

```
##           nprune=12)
```

```
##
```

```
##               coefficients
```

```
## (Intercept)      9748.1791
```

```
## h(Apps-3646)      0.4632
```

```
## h(2279-Accept)    -1.8972
```

```
## h(913-Enroll)     6.3690
```

```
## h(Enroll-913)    -1.9849
```

```
## h(1363-F.Undergrad) -1.9452
```

```
## h(5895-Room.Board) -0.6939
```

```
## h(1230-Personal)  0.8542
```

```
## h(perc.alumni-6)  25.4428
```

```
## h(Expend-6864)    0.7506
```

```
## h(Expend-15387)   -0.7703
```

```
## h(83-Grad.Rate)   -28.1240
```

```
##
```

```
## Selected 12 of 22 terms, and 9 of 16 predictors (nprune=12)
```

```
## Termination condition: RSq changed by less than 0.001 at 22 terms
```

```
## Importance: Expend, Grad.Rate, Room.Board, Accept, Enroll, F.Undergrad, ...
```

```
## Number of terms at each degree of interaction: 1 11 (additive model)
```

```
## GCV 2729781    RSS 1111485831    GRSq 0.8134662    RSq 0.8312207
```

```
coef(mars.fit$finalModel)
```

```
##      (Intercept)      h(Expend-15387)      h(83-Grad.Rate)      h(5895-Room.Board)
```

```
##      9748.1790945      -0.7702870      -28.1240387      -0.6939388
```

```
## h(1363-F.Undergrad)      h(1230-Personal)      h(Apps-3646)      h(Enroll-913)
```

```
##      -1.9451538      0.8542158      0.4631943      -1.9848536
```


```
##      h(913-Enroll)      h(2279-Accept)      h(perc.alumni-6)      h(Expend-6864)
```

```
##      6.3690181      -1.8971697      25.4427652      0.7506273
```



```
# partial dependence plot on arbitrary predictor Grad.Rate
p1 <- pdp::partial(mars.fit, pred.var = c("Grad.Rate"), grid.resolution = 10) |>
  autoplot()
```

p1



```
# test error
pred.mars <- predict(mars.fit, newdata = testing_data)

test.error.mars <- mean((pred.mars - y2)^2)
```

The regression function for the MARS model is $f(\text{Outstate}) = 9748.1791 + 0.4632 h(\text{Apps}-3646) - 1.8972 h(2279-\text{Accept}) + 6.3690 h(913-\text{Enroll}) - 1.9849 h(\text{Enroll}-913) - 1.9452 h(1363-\text{F.Undergrad}) - 0.6939 h(5895-\text{Room.Board}) + 0.8542 h(1230-\text{Personal}) + 25.4428 h(\text{perc.alumni}-6) + 0.7506 h(\text{Expend}-6864) - 0.7703 h(\text{Expend}-15387) - 28.1240 h(83-\text{Grad.Rate})$. The test error is 3.4360573×10^6 .

1c) Construct a generalized additive model (GAM) to predict the response variable. Does your GAM model include all the predictors? For the nonlinear terms included in your model, generate plots to visualize these relationships and discuss your observations. Report the test error.


```
set.seed(2)


# fit a GAM model using 10-fold cross-validation
gam.fit <- train(x, y,
  method = "gam",
  tuneGrid = data.frame(method = "GCV.Cp", select = c(TRUE, FALSE)),
  trControl = ctrl)


# for non-linear terms, generate plots to visualize relationships
gam.fit$finalModel


##
## Family: gaussian
## Link function: identity
##
## Formula:
## .outcome ~ s(perc.alumni) + s(Terminal) + s(Books) + s(PhD) +
##      s(Grad.Rate) + s(Top10perc) + s(Top25perc) + s(S.F.Ratio) +
##      s(Personal) + s(P.Undergrad) + s(Room.Board) + s(Enroll) +
##      s(Accept) + s(F.Undergrad) + s(Apps) + s(Expend)
##
## Estimated degrees of freedom:
## 2.17 1.00 1.00 2.32 4.49 7.59 1.00
## 3.74 1.00 1.00 2.50 1.00 2.84 6.17
## 3.89 7.39 total = 50.1
##
## GCV score: 2689385


plot(gam.fit$finalModel)
```













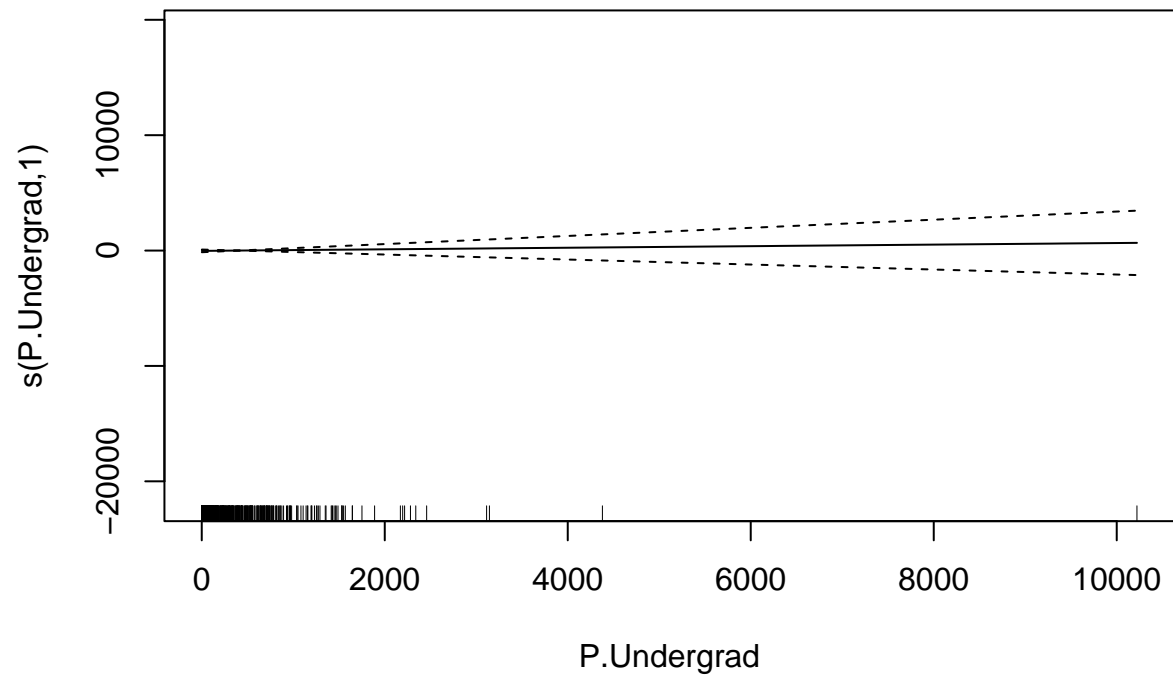






























```
# report test error
pred.gam <- predict(gam.fit, newdata = testing_data)

test.error.gam <- mean((pred.gam - y2)^2)
```

The best fit GAM model does include all predictors, as the selection method is **FALSE**. From the final models plots, we see that **Terminal**, **Books**, **Top25Perc**, **Personal**, and **Enroll** appear to be linear terms as their $df = 1$, but the all the predictors in the model have an “s”, indicating the smoothing function was used and these terms are interpreted as non-linear. We can observe that some predictors are more non-linear than others as the degrees of freedom increase. For example, **Top10perc** has more wave in the graph compared to **S.F.Ratio**. The test error is 3.4345234×10^6 .

1d) In this dataset, would you favor a MARS model over a linear model for predicting out-of-state tuition? If so, why? More broadly, in general applications, do you consider a MARS model to be superior to a linear model? Please share your reasoning.

```
set.seed(2)

# fit a linear model using 10-fold cross-validation
lm.fit <- train(x, y,
               method = "lm",
               trControl = ctrl)

summary(lm.fit)
```


```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6611  -1248       -7    1349   9810
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  752.97152   963.39650    0.782  0.43489
## Apps         -0.04677    0.11802   -0.396  0.69207
## Accept         1.33042    0.20335    6.542 1.70e-10 ***
## Enroll        -3.83350    0.91337   -4.197 3.28e-05 ***
## Top10perc     30.21342   15.52858    1.946  0.05234 .
## Top25perc      0.22821   12.58307    0.018  0.98554
## F.Undergrad    0.06437    0.14400    0.447  0.65508
## P.Undergrad   -0.19878    0.16435   -1.210  0.22711
## Room.Board     0.87381    0.11536    7.575 2.18e-13 ***
## Books          0.36615    0.54397    0.673  0.50123
## Personal      -0.48636    0.15441   -3.150  0.00175 **
## PhD            6.58415   11.93575    0.552  0.58148
## Terminal      32.39712   13.16108    2.462  0.01422 *
## S.F.Ratio     -37.67851   31.82476   -1.184  0.23708
## perc.alumni    39.30520    9.53167    4.124 4.47e-05 ***
## Expend         0.16060    0.02711    5.923 6.42e-09 ***
## Grad.Rate     20.39531    7.25323    2.812  0.00515 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2000 on 435 degrees of freedom
## Multiple R-squared:  0.7358, Adjusted R-squared:  0.7261
## F-statistic: 75.73 on 16 and 435 DF, p-value: < 2.2e-16

# compare models
resamp <- resamples(list(lm = lm.fit, mars = mars.fit))
```


```
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: lm, mars
## Number of resamples: 10
##
## MAE
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max. NA's
## lm  1427.528 1578.173 1615.828 1608.826 1674.692 1752.230    0
## mars 1138.836 1239.576 1320.672 1321.704 1425.698 1497.191    0
##
## RMSE
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max. NA's
## lm  1770.494 1950.509 2030.364 2047.597 2114.502 2393.058    0
## mars 1377.340 1548.527 1712.939 1688.556 1831.993 1956.649    0
##
## Rsquared
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max. NA's
## lm  0.6577688 0.7142791 0.7249317 0.7204843 0.7408093 0.7704528    0
## mars 0.7384265 0.7672405 0.8068437 0.8053855 0.8427477 0.8636907    0
```

```
parallelplot(resamp, metric = "RMSE")
```



```
bwplot(resamp, metric = "RMSE")
```



From the resampling summary, I believe the best model is the MARS since it has the smallest mean and median RMSE value. I would prefer fitting a MARS model over a linear model mostly because of its flexibility. MARS models can capture non-linear relationships, can capture interactions between variables, and are generally more robust.